

Tenxo — Market Research Report

Decentralized GPU Compute Landscape (May 2026)

1. GPU Compute Market: The Big Picture

Market Size & Growth

Metric	Value	Source
Global GPU market (2026)	\$100-104B	Mordor Intelligence, Fortune BI
Global GPU market (2034)	\$642B (26.1% CAGR)	Fortune Business Insights
AI data center GPU market (2026)	\$27.9B	Precedence Research
AI data center GPU market (2030)	\$228B (13.7% CAGR)	MarketsandMarkets
NVIDIA data center revenue (FY2026E)	\$150B+	Silicon Analysts
Decentralized compute market (2026)	\$8.9B (25.7% CAGR)	TBRC
Decentralized compute market (2030)	\$22.5B	TBRC
Decentralized GPU annualized protocol revenue	\$200M+ (early 2026)	DeFiLlama, Dune Analytics
DePIN sector market cap	\$9.4B (Mar 2026, peak \$19.2B in Sep 2025)	CoinGecko

Key Observations

- NVIDIA dominates** ~80-90% of AI accelerators. Their data center revenue grew from \$15B (2022) to \$130B+ (2025E). Even as share declines to ~75% by 2026, absolute revenue keeps growing because the market expands faster than competitors can capture share.
- Inference > Training** — The AI inference market is the largest and fastest-growing segment. Most decentralized networks target inference, not training, because inference is highly parallelizable and latency-tolerant.
- The GPU shortage is structural** — AWS H100 clusters had 8-12 month waitlists in 2023-2024. Cloud providers raised AI compute prices 40-60% in 2025. McKinsey

estimates the global AI infrastructure market will surpass **\$700B annually by 2030**.

- 4. **Cost differential is massive** — Decentralized GPU networks offer **60-90% cost savings** vs AWS/Azure comparable instances. Akash settles at 80-90% below AWS for CPU; Render benchmarks ~70% below Azure ML compute.

2. The Decentralized GPU Competitor Landscape

Tier 1: The "Big Three" DePIN Compute Networks

Render Network (RENDER)

Dimension	Detail
Founded	2017 (Ethereum), migrated to Solana 2023-2024
Market Cap	~\$700M - \$1.5B (varies)
Core Focus	GPU rendering → expanded to AI inference
Revenue Model	Burn-Mint Equilibrium (95% of job spend burned)
2025 Revenue	Not publicly disclosed
Strengths	Longest track record, real enterprise customers (Pixar, HBO, Apple via OTOY), verified job results
Weaknesses	Permissioned onboarding (Foundation approval required), still rendering-heavy, opaque pricing
Security	No TEE/confidential compute in production. Relies on reputation-based node selection (Tier 1/2/3). Node operators must be approved by Foundation. Job verification is about output quality (render correctness), not privacy.
Crypto Mining Prevention	None publicly documented. Permissioned onboarding acts as a soft barrier — Foundation can deny known mining operations. However, no runtime detection exists. Nodes run arbitrary GPU workloads.
Notable	Launched Dispersed.com (Dec 2025) — AI compute subnet with 600+ models at \$1.75/hr. Node operators must be approved by the Render Foundation.

Akash Network (AKT)

Dimension	Detail
Founded	2018 (Cosmos SDK)
Market Cap	~\$88M - \$146M
Core Focus	General-purpose decentralized cloud (CPU + GPU + storage)
Revenue Model	Burn-Mint Equilibrium (activated Mar 2026) + take fee on deployment spend
2025 Revenue	\$3.15M (on-chain audited, 128% YoY growth). Q1 2026 ATH: \$5M
Strengths	Permissionless, open-source (Apache 2.0), verifiable revenue, Kubernetes-native, reverse auction drives lowest prices
Weaknesses	Supply inconsistency, no SLAs, general-purpose means not optimized for AI
Security	Most advanced in TEE. AEP-83/Kata Containers: confidential computing via Intel TDX + AMD SEV-SNP + NVIDIA H100 CC mode. Composite attestation (CPU + GPU). Still in roadmap/proposal stage (AEP-65, AEP-83). Not fully deployed in production yet.
Crypto Mining Prevention	None explicit. Relies on marketplace dynamics (providers compete on price; miners would be outbid by serious providers) + reputation system. No workload inspection. Permissionless means anyone can join.
Notable	Starcluster initiative — acquiring ~7,200 NVIDIA GB200 GPUs for enterprise-grade "planetary mesh." Planning to migrate off Cosmos SDK by late 2026.

io.net (IO)

Dimension	Detail
Founded	2022 (Solana)
Market Cap	~\$32M - \$400M+ (peak during growth cycle)
Core Focus	AI GPU aggregation (training + inference)
Revenue Model	Incentive Dynamic Engine (IDE) launching Q2 2026
2025 Revenue	~\$12.5M annualized (self-reported, partially verified)
Strengths	Largest GPU pool (claims 100K+ devices), AI-optimized clustering, low cost (90% below cloud), enterprise partnerships (Dell, KREA/Nike, UC Berkeley)
Weaknesses	Sybil attack history (1.8M fake GPUs in 2024), only ~6,720 daily verified active, self-reported revenue, verification gap
Security	No TEE support documented. Uses "proof-of-work" mechanism to verify GPUs are online (cryptographic verification tasks), but this is for availability verification, not privacy. No confidential compute.
Crypto Mining Prevention	None documented. In fact, io.net explicitly aggregates GPUs from crypto mining farms transitioning away from PoW mining. This is a feature, not a bug — mining farms are a major supply source. Their "proof-of-work" verification mechanism only checks that a GPU is present and responding, not what it's running.
Notable	Biggest hype-to-reality gap. P/Revenue ~2.6x (market is skeptical). Joined Dell partner program (Dec 2024).

Tier 2: Centralized GPU Marketplace Competitors

These are not "decentralized" in the blockchain sense — they are centralized platforms that aggregate GPU supply from datacenters and individuals. They are Tenxo's competition for developer mindshare and GPU supply.

RunPod

Dimension	Detail
Founded	2020 (US)
Business Model	Centralized GPU cloud — two tiers: Secure Cloud (T3/T4 datacenters, SLA) + Community Cloud (peer-to-peer hosts, no SLA) + Serverless (scale-to-zero inference)
ARR (2026)	~ \$120M (90% YoY growth, only \$20M raised — extremely capital efficient)
Funding	Intel Capital, Dell Technologies Capital (strategic). A16z Speedrun default GPU infra partner
GPU Catalog	18+ GPU classes: RTX 4090/5090 → H100/H200/B200. Consumer to enterprise
Pricing	RTX 4090: \$0.34/hr (Community), H100: \$1.99-2.79/hr, A100: \$1.19/hr. Per-second billing
Differentiator	Best balance of price + usability. Docker-native, sub-200ms cold starts on Serverless, template ecosystem
Security	SOC 2 Type II (Oct 2025), HIPAA, GDPR. No TEE/confidential compute. Workloads run in plaintext in Docker containers. Community Cloud hosts are not audited — data could be observed by host operator.
Crypto Mining Prevention	None for Community Cloud. Community Cloud is peer-to-peer (like Vast.ai) — hosts are vetted but not monitored for workload type. Secure Cloud is datacenter-only with contractual prohibitions. The platform has no runtime workload detection.
Multi-Node	Instant Clusters (up to 8 GPUs on one host). No multi-host InfiniBand clustering
Strengths	Massive GPU selection, excellent developer experience, per-second billing, SOC 2 certified, template library (1-click deploy popular models), serverless auto-scaling
Weaknesses	Community Cloud availability is spiky (capacity exhausts in peak hours). No multi-node training at scale. Community Cloud reliability varies. No privacy guarantees for workloads.
Strategic Note	RunPod is Tenxo's most relevant "traditional" competitor because they bridge the gap between hyperscaler and p2p. They are the default choice for AI devs who want cheap GPUs without dealing with crypto. Their weakness: no privacy. Every workload runs in plaintext visible to the host.

Vast.ai

Dimension	Detail
Founded	2018 (Los Angeles)
Business Model	Peer-to-peer GPU marketplace. Individual hosts (gamers, ex-miners, small colo) list GPUs. Vast takes ~25-30% commission. Pure marketplace — no owned infrastructure
ARR (2026)	~\$2.2M (Vast's platform revenue; underlying GMV is much higher since this is commission-only)
Funding	~\$4M from DRW Holdings, Nazare. Largely bootstrapped
GPU Catalog	68+ GPU classes — largest selection in market. RTX 3060 → B300 SXM. 17,000+ GPUs across 1,400+ providers in 500+ locations
Pricing	Cheapest in market. RTX 4090: ~\$0.20-0.35/hr (vs RunPod \$0.34), A100 80GB: ~\$0.67-1.00/hr, H100: ~\$0.90-1.87/hr. Per-second billing. Lower than all competitors by 30-60%
Differentiator	Absolute lowest price. Largest consumer GPU selection. Marketplace dynamics drive prices below fixed providers
Security	SOC 2 Type II, HIPAA, GDPR (on Secure Cloud tier). No TEE. "Verified hosts" have reliability scores and benchmarks. Fundamental trust problem: you're renting someone's personal machine — host has physical access, your data runs in plaintext in Docker.
Crypto Mining Prevention	None. In fact, ex-crypto-miners are a major supply source. The platform embraces this. No workload inspection. Host reliability scores help but don't detect mining.
Multi-Node	Single-host only. No native multi-host clustering. Cannot do distributed training across machines
Strengths	Lowest prices in the market by significant margin. Widest GPU selection (especially consumer GPUs). Per-second billing. Spot/interruptible bidding for 50%+ discounts
Weaknesses	Reliability varies wildly by host (residential internet drops kill jobs). No formal support (Discord/email only). Complex UI. No SLA. Security/compliance is the weakest of any major provider — your code runs on a stranger's machine in plaintext.
Strategic Note	Vast.ai demonstrates that there IS massive untapped consumer GPU supply — their 17,000+ GPUs from 1,400+ hosts proves the thesis. But their model has no privacy at all. Tenxo can go after the same supply with the differentiator of "your code stays encrypted." Vast.ai also shows that pure marketplace models suffer from reliability problems that turn off serious AI teams.

Comparison: RunPod vs Vast.ai vs Tenxo

Feature	Tenxo	RunPod	Vast.ai
Category	Decentralized (ZK)	Centralized (hybrid)	Centralized (P2P)
Privacy	☐ TEE + ZK routing	× Plaintext Docker	× Plaintext Docker
Consumer GPUs	☐ (RTX 4090, etc.)	☐ (Community tier)	☐ (68+ classes)
Enterprise GPUs	Planned	☐ (Secure Cloud H100/B200)	☐ (via verified hosts)
Fiat Payments	☐	☐	☐
Pricing (RTX 4090)	\$0.15/hr	\$0.34/hr (Community)	\$0.20-0.35/hr
Pricing (H100)	\$0.75/hr	\$1.99-2.79/hr	\$0.90-1.87/hr
Reliability	TEE-verified	High (Secure) / Variable (Community)	Variable (by host)
Security Cert	None yet	SOC 2, HIPAA, GDPR	SOC 2, HIPAA, GDPR
Multi-Node	Planned	Up to 8 GPU	None
Revenue	Pre-revenue	~\$120M ARR	~\$2.2M ARR (commission)
Mining Prevention	None yet	None	None (embraces miners)

Tier 3: Specialized / Emerging Players

Aethir

Dimension	Detail
Core Focus	Enterprise GPU for cloud gaming + AI training
Revenue	~\$150M annualized (Q3 2025: \$39.8M quarterly). Largest in sector.
Key Strength	Enterprise SLAs, 150+ paying enterprise clients
Security	No public TEE/confidential compute documentation. Enterprise focus suggests contractual security, not cryptographic.
Crypto Mining	Not addressed publicly. Enterprise SLAs would prohibit mining, but enforcement is contractual.
Platform	Arbitrum-based

Gensyn

Dimension	Detail
Core Focus	Decentralized AI training (not inference) with cryptographic verification
Key Innovation	"Proof-of-Learning" — cryptographically verify that training actually happened
Backing	a16z (\$43M Series A)
Status	Token launched 2026; market cap ~\$71M circulating, \$546M FDV
Security	Groundbreaking approach: uses cryptographic proofs to verify computation correctness, not hardware TEEs. Different trust model — you can prove training was done correctly without trusting hardware.
Crypto Mining	Proof-of-Learning would theoretically detect mining since the cryptographic proof wouldn't match expected training computation. But this is unproven at scale.

Hyperbolic

Dimension	Detail
Core Focus	AI inference at 75% below cloud
Key Metric	Powers 100,000+ developers
Architecture	Hyper-dOS — decentralized operating system for GPU orchestration
Security	Planning "cryptographic verification" but not yet deployed

Nosana (NOS)

Dimension	Detail
Core Focus	AI inference at the edge (low-latency)
Platform	Solana-based
Niche	Latency-sensitive AI applications
Security	No public TEE documentation

Flux / FluxCloud

Dimension	Detail
Core Focus	Full-stack decentralized cloud (CPU + GPU + storage + apps)
Differentiator	"Platform supply" vs just "compute supply" — provides runtime environments, not just raw GPU
Security	No TEE documented. Competes on completeness of platform, not privacy guarantees.

Cocoon (TON-based)

Dimension	Detail
Core Focus	Decentralized AI inference with TEE privacy
Architecture	RA-TLS encrypted connections, TEE attestation (Intel TDX), on-chain registries
Security	TEE-first design. Uses Intel TDX for both proxy and worker nodes. On-chain image hash registry for verification. GPU verification on boot (fail-closed).
Crypto Mining	Not documented. TEE-based design would technically make it hard to mine (mining software might not run in TDX), but this is incidental.

Cerumbra

Dimension	Detail
Core Focus	End-to-end encrypted browser-to-TEE GPU inference
Architecture	Browser ECDH → AES-GCM → NVIDIA H100 TEE / Blackwell TEE
Status	Pre-production — waiting for H100 TEE hardware availability
Security	Deepest TEE integration. Protocol-level encryption from browser to GPU. Remote attestation with full certificate chain verification.
Crypto Mining	Natural prevention: mining workloads can't run in NVIDIA H100 TEE (designed for inference only).

Phala Network

Dimension	Detail
Core Focus	TEE-based confidential computing for blockchain + AI
Key Innovation	"Private ML SDK" — turn Docker workflows into TEE-protected deployments
Security	Most advanced TEE implementation. Intel SGX → TDX + AMD SEV-SNP + NVIDIA H100 GPU TEE. On-chain attestation anchoring. Composite CPU+GPU attestation.
Crypto Mining	Not explicitly addressed. TEE architecture inherently restricts what can run (signed measurements). Mining software wouldn't match expected measurements.

Competitive Matrix Summary (Full Market)

Feature	Tenxo	RunPod	Vast.ai	Render	Akash	io.net	Aethir
Category	Decentralized ZK	Centralized hybrid	Centralized P2P	DePIN	DePIN	DePIN	DePIN
Consumer GPUs	☐	☐ (Community)	☐ (68+ classes)	☐	☐	☐	×
Enterprise GPUs	Planned	☐ (Secure Cloud)	☐ (verified hosts)	☐	☐ (Starcluster)	☐	☐
Fiat Payments	☐	☐	☐	×	×	×	×
TEE / Privacy	☐ SEV-SNP	×	×	×	☐ roadmap	×	×
ZK Routing	☐	×	×	×	×	×	×
Revenue (2025)	Pre-revenue	~\$120M ARR	~\$2.2M ARR*	Undisclosed	\$3.15M	\$12.5M (claimed)	~\$150M
Security Cert	None yet	SOC 2, HIPAA, GDPR	SOC 2, HIPAA, GDPR	None	None	None	None
Multi-Node	Planned	Up to 8 GPU	None	☐	☐ (Starcluster)	☐ (clusters)	☐
Mining Prev.	None yet	None	None (embraces)	Permissioned gate	Marketplace	Feature (farms)	Contractual

* Vast.ai ARR is commission-only; underlying GMV is much higher.

What this tells you: Tenxo is the only platform offering privacy + fiat + consumer GPUs in one package. RunPod and Vast.ai have the revenue and the supply but zero privacy guarantees. The DePIN competitors have blockchain but no privacy (except Akash on roadmap). Tenxo's window is the gap between "developers want cheap compute" and "developers need their code to not leak."

3. Security & Privacy: How Competitors Handle It

The Spectrum of Approaches

No Privacy	Full Privacy
v	v

io.net	Render	Akash (planned)	Cocoon	Tenxo	Cerumbra	Phala
(none)	(reputation)	(TEE roadmap)	(TDX)	(SEV-SNP)	(H100 TEE)	(Multi-TEE)

Key Findings

1. **Most competitors have NO privacy guarantees.** io.net, Render, and Aethir run workloads in plaintext on untrusted machines. Privacy is not part of their value prop.
2. **Akash is moving toward TEE** but it's still on the roadmap (AEP-65, AEP-83). Kata Containers with Intel TDX + AMD SEV-SNP + NVIDIA H100 CC-mode. Not deployed in production yet.
3. **TEE-first projects exist but are pre-production.** Cocoon, Cerumbra, and Phala have the right architecture but lack network effects, supply, or production readiness. Phala is the most advanced but uses Intel SGX (limited memory, 128MB enclaves) which severely limits AI workloads.
4. **Gensyn takes a unique approach** — cryptographic "Proof-of-Learning" instead of hardware TEEs. This is academically interesting but unproven at scale. No production benchmarks exist.
5. **TEE hardware availability is a bottleneck** — NVIDIA H100/H200 confidential compute mode requires specific firmware and is only available on the newest hardware. AMD SEV-SNP is more widely available on EPYC CPUs. Tenxo's choice of AMD SEV-SNP is pragmatic — more existing infrastructure supports it.

Tenxo's Position

Tenxo is **the only production-ready platform combining all three**: - ☐ TEE attestation (AMD SEV-SNP, working today) - ☐ Zero-knowledge routing (matchmaker never sees plaintext) - ☐ Consumer GPU support (RTX 4090, etc.) - ☐ Fiat payments (Razorpay, no crypto)

vs Akash (TEE on roadmap, crypto-only) vs Render (no TEE, crypto-only, permissioned) vs io.net (no TEE, crypto-only, supply quality issues)

4. Crypto Mining Abuse: How Competitors Handle It

This is the question you asked earlier. Here's the honest state of the industry:

Nobody Has a Good Answer

Competitor	Approach	Effectiveness
io.net	Mining farms are a feature. They explicitly aggregate GPUs from crypto mining operations. Their "proof-of-work" verification only checks GPU availability.	Not prevention — supply source
Render	Permissioned onboarding — Foundation must approve nodes. Soft barrier but runtime detection doesn't exist.	Low (rejects known miners only)
Akash	Marketplace dynamics + reputation. Permissionless — anyone can join. No workload inspection.	Low (miners just need to bid competitively)
Aethir	Enterprise SLAs — contractual prohibition. Enforcement is legal, not technical.	Moderate (contractual only)
Gensyn	Proof-of-Learning would theoretically detect non-training workloads.	Untested at scale
Phala / Cocoon / Cerumbra	TEE-restricted execution — only signed/measured workloads can run.	High (mining software wouldn't match expected measurements)
Tenxo	Currently none (same as most competitors)	N/A

What Actually Works

1. **TEE-restricted execution** (Phala model) — Only workloads whose measurements match expected values can run. Mining software would fail attestation. **Downside:** requires curated image registry, limits flexibility.
2. **GPU profiling** (proposed for Tenxo) — Monitor utilization patterns. Mining: flat 100%, predictable memory. ML: bursty, variable memory, I/O wait. **Cheap, effective for naive miners.**
3. **Financial barriers** (Tenxo's current approach) — Prepaid credits (\$5 minimum). Wallets are authenticated. Mining may not be profitable at \$0.15/hr RTX 4090 pricing vs cloud GPU mining alternatives.
4. **Output heuristics** — Check result sizes, file types for known mining artifacts (wallet files, DAG files). **Metadata only, no plaintext exposure.**
5. **Reputation & staking** — Providers stake tokens. If mining detected, stake is slashed. **Requires tokenomics (Tenxo is fiat-only).**

Honest Assessment

No decentralized GPU network has solved this problem. The fundamental tension is: - Zero-knowledge means the platform can't inspect workloads - Preventing abuse requires some inspection

The industry consensus is that **pricing arbitrage + financial barriers + reputation systems** are sufficient for MVP. Mining at decentralized GPU prices (\$0.15-0.50/hr for consumer GPUs) is less profitable than dedicated mining on ASICs or cloud GPU mining at scale. The risk is real but manageable.

Tenxo's proposed layered approach (profiling → heuristics → staking) would be **more comprehensive than most competitors** while preserving the zero-knowledge architecture.

5. Strategic Insights for Tenxo

What Competitors Get Wrong

1. **Crypto-only payments** — Every major competitor forces users to hold volatile tokens. This limits adoption to crypto-native users. Tenxo's fiat-only approach (Razorpay) addresses the **real** market: AI developers who don't want to touch crypto.
2. **No privacy guarantees** — io.net, Render, and Akash run code in plaintext. The #1 question from enterprise AI teams is "will my model weights leak?" Tenxo can answer "no, they're encrypted end-to-end and run in a TEE." This is a **massive differentiator**.
3. **Permissioned or permissionless extremes** — Render is too centralized (Foundation approval). Akash is too chaotic (anyone can join, inconsistent quality). Tenxo's model (permissionless for supply, TEE-verified execution) offers the best balance.
4. **GPU supply quality** — io.net's Sybil attack showed the problem with unverified supply. Tenxo's TEE attestation + zero-knowledge routing means **providers can't fake their hardware** and **can't steal your data**.

What Tenxo Should Watch

1. **NVIDIA's response** — If NVIDIA launches a decentralized compute product (unlikely but possible), it could reshape the market given their ~80% market share.
2. **Akash's TEE rollout** — If Akash successfully deploys Kata Containers + TEE, they eliminate Tenxo's privacy advantage. But they'd still require crypto tokens.
3. **Enterprise adoption timing** — The market isn't asking for decentralized compute yet. AI startups are desperate for cheap compute and will sacrifice privacy for cost. Tenxo needs to compete on price first, privacy second.

4. **Supply-side chicken-and-egg** — Without GPU providers, there's no compute to sell. Without compute, there are no developers. The installer script + TEE attestation makes joining as a provider easy, but Tenxo needs a critical mass of GPUs to be useful.

Tenxo's Unique Value Proposition

Property	Tenxo	Competitors
Privacy	☑	× (mostly)
Consumer GPUs	☑	☑ (some)
Fiat payments	☑	× (all crypto)
TEE verified	☑	× (mostly not)
Open source	Partial	☑/× mixed
Production	☑ MVP	☑ (some)

The **unbeatable combination** is: **Pay with a credit card, your code stays private, it runs on cheap consumer GPUs, and you can verify it happened correctly.** No competitor offers this package today.

Research compiled May 2026. Sources: CoinGecko, DeFiLlama, Messari, Dune Analytics, project documentation, market research reports from Fortune BI, Mordor Intelligence, MarketsandMarkets, Precedence Research, TBRC, Grand View Research.

References & Sources

Market Size & GPU Industry

1. **Decentralized Compute Global Market Report 2026** — TBRC / GII Research <https://www.giiresearch.com/report/tbrc1987682-decentralized-compute-global-market-report.html>
2. **2026 Real Progress and Investment Opportunities in Decentralized AI Computing Networks** — Bee Network <https://www.bee.com/72099.html>
3. **AI Compute Demand Is Outpacing Supply, And Crypto Networks Are Stepping In** — Yellow.com <https://yellow.com/research/ai-compute-demand-crypto-gpu-networks-gap-2026>

4. **Decentralized GPU Networks: The Flow of AI Compute Dollars** — Alinvest <https://www.ainvest.com/news/decentralized-gpu-networks-flow-ai-compute-dollars-2602/>
5. **Decentralized GPU Networks 2026: How DePIN is Challenging AWS for the 00B AI Compute Market** — Blockeden <https://blockeden.xyz/blog/2026/02/07/decentralized-gpu-networks-2026/>
6. **Decentralized GPU Marketplaces vs Big Tech Cloud** — Perspective Labs <https://perspectivelabs.org/decentralized-gpu-vs-big-tech-cloud/>
7. **Decentralized Compute Market Research Report 2033** — Growth Market Reports <https://growthmarketreports.com/report/decentralized-compute-market>
8. **AI Data Center GPUs Global Market Report 2026** — TBRC <https://www.giiresearch.com/report/tbrc1987574-artificial-intelligence-ai-data-center-graphics.html>
9. **Data Center GPU Market Size & Share Report 2033** — Grand View Research <https://www.grandviewresearch.com/industry-analysis/data-center-gpu-market-report>
10. **Data Center GPU Market Size, Share & Trends 2025-2030** — MarketsandMarkets <https://www.marketsandmarkets.com/Market-Reports/data-center-gpu-market-18997435.html>
11. **NVIDIA AI GPU Market Share 2026: ~80% of AI Accelerators** — Silicon Analysts <https://siliconanalysts.com/analysis/nvidia-ai-accelerator-market-share-2024-2026>
12. **Graphics Processing Unit Market Size & Forecast 2033** — Persistence Market Research <https://www.persistencemarketresearch.com/market-research/graphics-processing-unit-market.asp>
13. **Data Center GPU Market Size to Reach USD 226.87 Billion By 2035** — Precedence Research <https://www.precedenceresearch.com/data-center-gpu-market>
14. **GPU Market Size, Trends, Share & Industry Report 2031** — Mordor Intelligence <https://www.mordorintelligence.com/industry-reports/graphics-processing-unit-market>
15. **Graphic Processing Unit (GPU) Market Size, Share & Forecast 2034** — Fortune Business Insights <https://www.fortunebusinessinsights.com/graphic-processing-unit-gpu-market-109831>

Competitor Analysis (DePIN)

1. **Render vs io.net vs Akash: DePIN Compute Networks Comparison** — Gate Learn <https://www.gate.com/learn/articles/render-vs-ionet-vs-akash-depin-comparison>
2. **Render vs Akash vs io.net 2026: Revenue, Burns & Tokenomics** — Own Your Mind <https://ownyourmind.ai/tokenomics/render-vs-akash-vs-ionet/>
3. **Decentralized AI Compute: GPU Networks Powering the AI Revolution** — Whale Factor <https://www.whalefactor.com/ai/compute>
4. **Render Network GPU Economy Hits .2B** — Nonce Media <https://www.noncemedial.com/render-network-gpu-economy-ai-compute-2026/>
5. **The Decentralized AI Compute Landscape: Gensyn, Render, and Akash** — Gate.com <https://www.gate.com/blog/102500/decentralized-ai-compute-gensyn-render-akash-competition-landscape-web3-ai-infrastructure-crypto-compute-market-analysis>
6. **What Is Render Network? Decentralized GPU DePIN Guide 2026** — DEXTools <https://www.dextools.io/tutorials/what-is-render-network-rndr-decentralized-gpu-rendering-depin-guide-2026>
7. **RENDER vs AKT: Which AI Compute Token Has the Stronger Case?** — MKN Crypto News <https://news.mkncrypto.com/render-vs-akt-which-ai-compute-token-has-the-stronger-case/>
8. **Flux Against the World: State of Compute Networks 2026** — RunOnFlux <https://runonflux.com/flux-against-the-world-state-of-compute-networks-2026/>
9. **Decentralized GPU & Compute Networks Guide 2026** — degen0x <https://degen0x.com/learn/decentralized-gpu-compute-guide-2026/>
10. **How Decentralized GPU Marketplaces Like Akash and Render Solve the AI Compute Crisis** — FinanceFeeds <https://financefeeds.com/how-decentralized-gpu-marketplaces-like-akash-and-render-solve-the-ai-compute-crisis/>
11. **Render Network Review** — Own Your Mind <https://ownyourmind.ai/projects/render/>
12. **What Is Akash Network? DePIN Cloud & AKT Guide 2026** — DEXTools <https://www.dextools.io/tutorials/what-is-akash-network-depin-cloud-compute-akt-guide-2026>
13. **DePIN's Revenue Reckoning: How Akash, io.net, and Aethir Are Replacing Token Mining** — Blockeden <https://blockeden.xyz/blog/2026/03/12/depin-compute-revenue-pivot-akash-ionet-aethir/>

Security, TEE & Confidential Computing

1. **Confidential Compute via Kata Containers** — Akash AEP-83 <https://akash.network/roadmap/aep-83/>
2. **Confidential Computing** — Akash AEP-65 <https://akash.network/roadmap/aep-65/>
3. **Privacy-Preserving Decentralized AI with Confidential Computing** — arXiv <https://arxiv.org/html/2410.13752v1>
4. **GPU TEE Deep Dive: Securing AI at the Hardware Layer** — Phala Network <https://phala.com/posts/Phala-GPU-TEE-Deep-Dive>
5. **Beyond SGX: Embracing GPU TEE for Decentralized AI** — Phala Network <https://phala.com/posts/beyond-sgx-embracing-gpu-tee-for-decentralized-ai-dagi>
6. **Cocoon Architecture — Confidential Compute Open Network** <https://cocoon.org/architecture>
7. **Cerumbra — End-to-End Encrypted GPU Conversations** <https://cerumbra.org/>
8. **dstack Security Model** — GitHub <https://github.com/Dstack-TEE/dstack/blob/master/docs/security/security-model.md>

Competitor Analysis (RunPod & Vast.ai)

1. **RunPod vs Lambda Labs vs Vast.ai: GPU Rental Compare 2026** — Dmytro Klymentiev <https://klymentiev.com/blog/runpod-vs-lambda-vs-vast>
2. **Best Cloud GPU Providers 2026: RunPod vs Lambda vs Vast.ai** — TechPlained <https://www.techplained.com/best-gpu-cloud-ai-training>
3. **RunPod Review 2026: The Developer-First GPU Cloud at 20M ARR** — jimmy-research <https://jimmyresearch.com/entities/runpod/>
4. **Runpod Reliability Deep-Dive — 20M ARR, Triple Certification** — Endplan <https://endplan.ai/reports/cloud-gpu/runpod-trust-report-2026-en>
5. **RunPod vs Vast.ai: We Tested Both — Here's the Winner (2026)** — Valebyte <https://valebyte.com/en/guides/runpod-vs-vastai-real-llm-inference-benchmarks-cost-analysis/>
6. **Runpod vs. Vast.ai: A Deep Dive into GPU Cloud Platforms** — DEV Community <https://dev.to/therabbithole/runpod-vs-vastai-a-deep-dive-into-gpu-cloud-platforms-for-aiml-10ga>
7. **Vast.ai Review 2026: The Cheapest GPU Cloud?** — BestGPUCloud <https://www.bestgpucloud.com/en/blog/vastai-review-2026>
8. **Vast.ai Review 2026: Pricing & Ratings** — VPSRated <https://www.vpsrated.com/provider/vast-ai>

9. **Vast.ai review** — Pintoed AI <https://pintoedai.com/reviews/vast>
 10. **Vast.ai Cloud GPU Review & Pricing — Updated May 2026** — CloudGPUPrice <https://cloudgpuprice.com/provider/vast-ai/>
 11. **Vast.ai Review 2026: Pricing, Reliability & Alternatives** — GPUUnex Blog <https://www.gpunex.com/blog/vast-ai-review-2026/>
 12. **Vast.ai — jimmy-research** <https://jimmyresearch.com/entities/vast-ai/>
 13. **Vast.ai Reviewed: Marketplace GPU Compute in 2026** — AI Chat Daily <https://www.aichatdaily.com/tools/vast-ai>
 14. **Genesis Cloud vs RunPod vs Vast.ai: Which is Best in 2026?** — GPUCloudList <https://www.gpucloudlist.com/en/blog/genesis-cloud-vs-runpod-vs-vastai>
-

All sources accessed May 2026. Prices and metrics are as of the publication dates noted in each source. Market figures may vary across sources due to differing methodologies and scope definitions.